

```
=====
```

Title: 2.2 Exercises

Author: Chad Wood

Date: 8 Dec 2021

Modified By: Chad Wood

Description: This program demonstrates the use of scatterplots and linear regression lines

```
=====
```

Scatterplots and Linear Regression

For this assignment, you will be using the data set car data.csv. This data set is a modified version of the data set imports 85.data from <https://archive.ics.uci.edu/ml/datasets/automobile>. There are two columns in this data set. One is the weight of the vehicle in pounds, and the other is the highway miles per gallon.

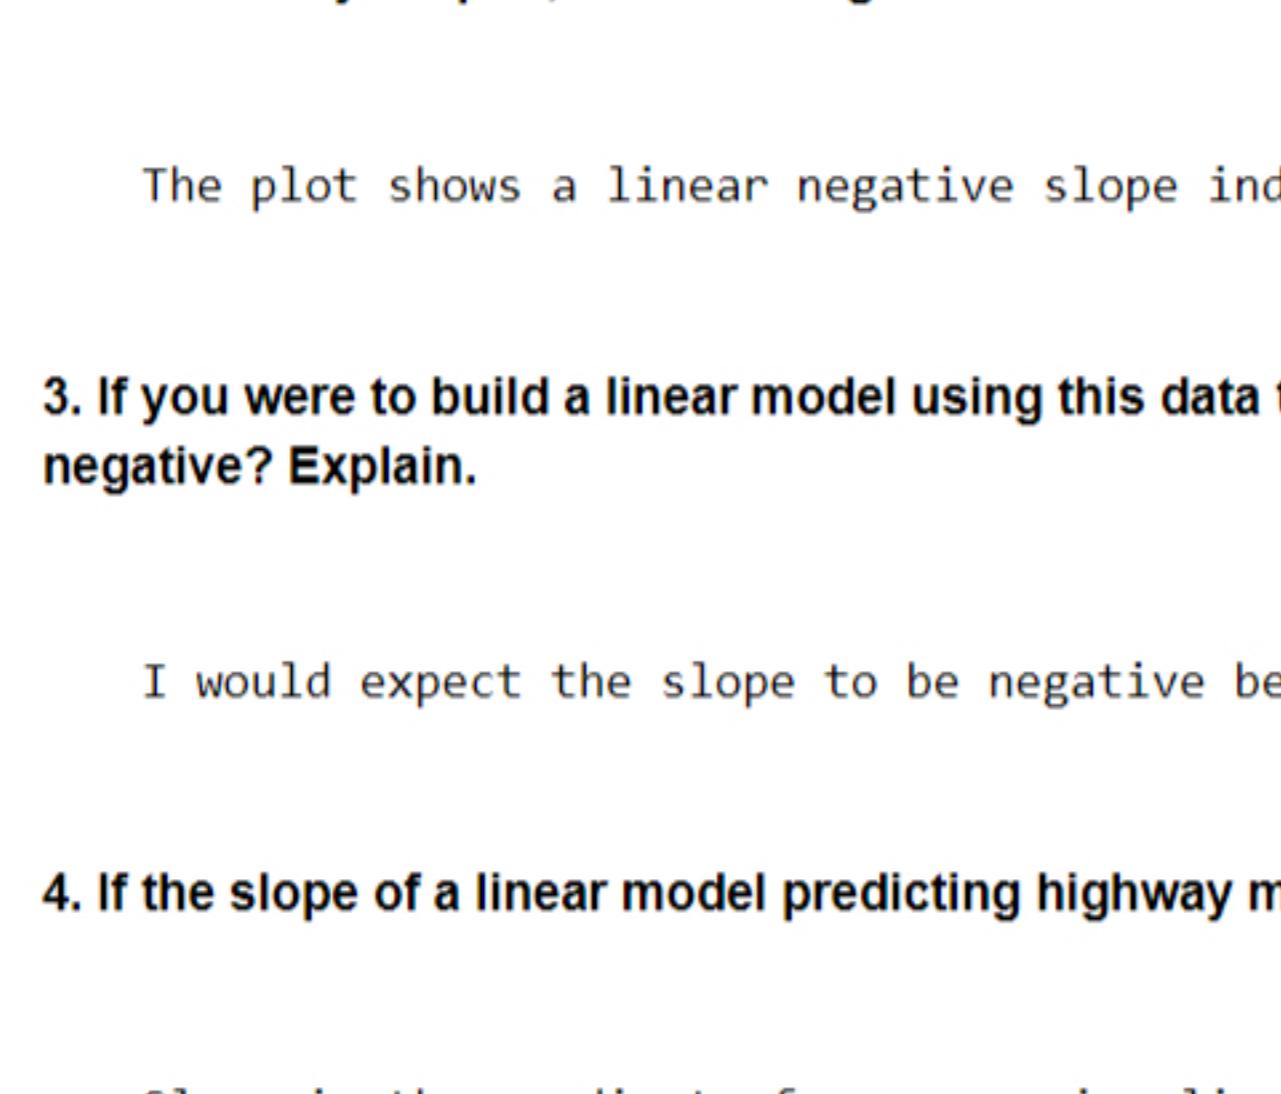
Setup

```
In [12]: import pandas as pd  
import matplotlib.pyplot as plt  
  
cars = pd.read_csv('week2data\car_data.csv')  
houses = pd.read_csv('week2data\housing_data.csv')
```

1. Make a scatterplot of the highway miles per gallon (y-axis) versus the weight (x-axis)

```
In [14]: # plt.scatter(x, y)  
plt.scatter(cars.weight, cars.hwy_mpg)
```

```
Out[14]: <matplotlib.collections.PathCollection at 0x1a43f0d8a00>
```



2. Based on your plot, what is the general trend of how highway miles per gallon varies with the weight?

The plot shows a linear negative slope indicating that as weight goes up, mpg goes down.

3. If you were to build a linear model using this data to predict highway miles per gallon from weight, would you expect the slope to be positive or negative? Explain.

I would expect the slope to be negative because it descends as the explanatory variable decreases.

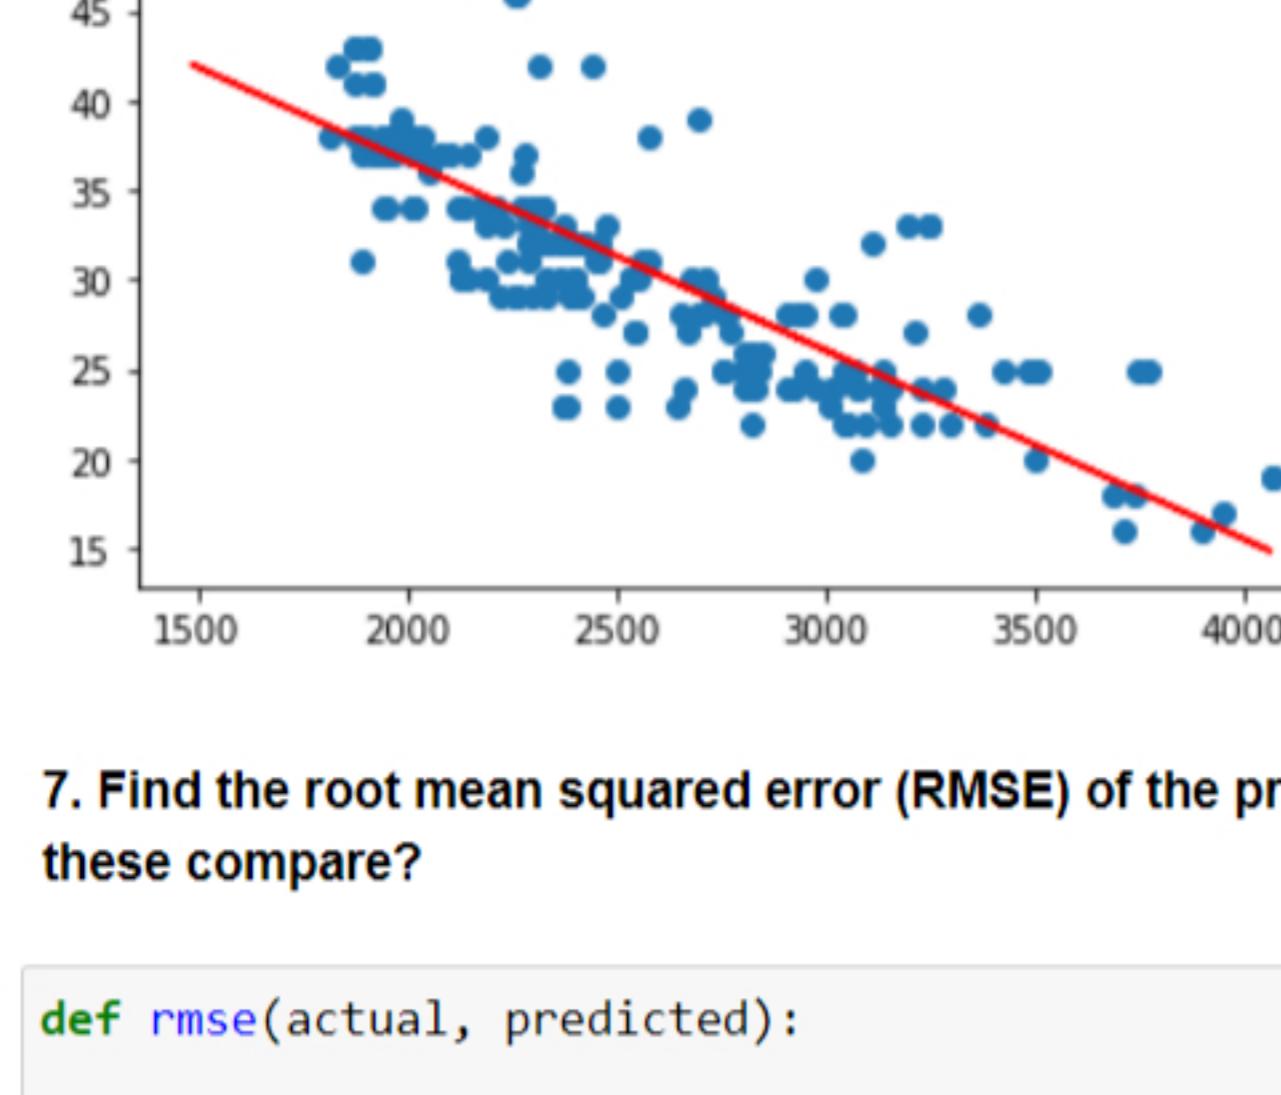
4. If the slope of a linear model predicting highway miles per gallon from the weight, interpret the meaning of the slope being -0.05.

Slope is the gradient of a regression line, and a slope of -0.05 would equate to a change in y of -0.05 for every change in x of 1

5. Write code to add a line to the graph you made in problem (1). Adjust the slope and y-intercept of this line until you think you have found the line that best fits the data. Record the slope and y-intercept.

```
In [31]: x, y = cars.weight, cars.hwy_mpg  
  
plt.scatter(x,y)  
  
# Adds regression line, m: -0.01; b: 56  
plt.plot(x, -0.01*x + 56, color='red')
```

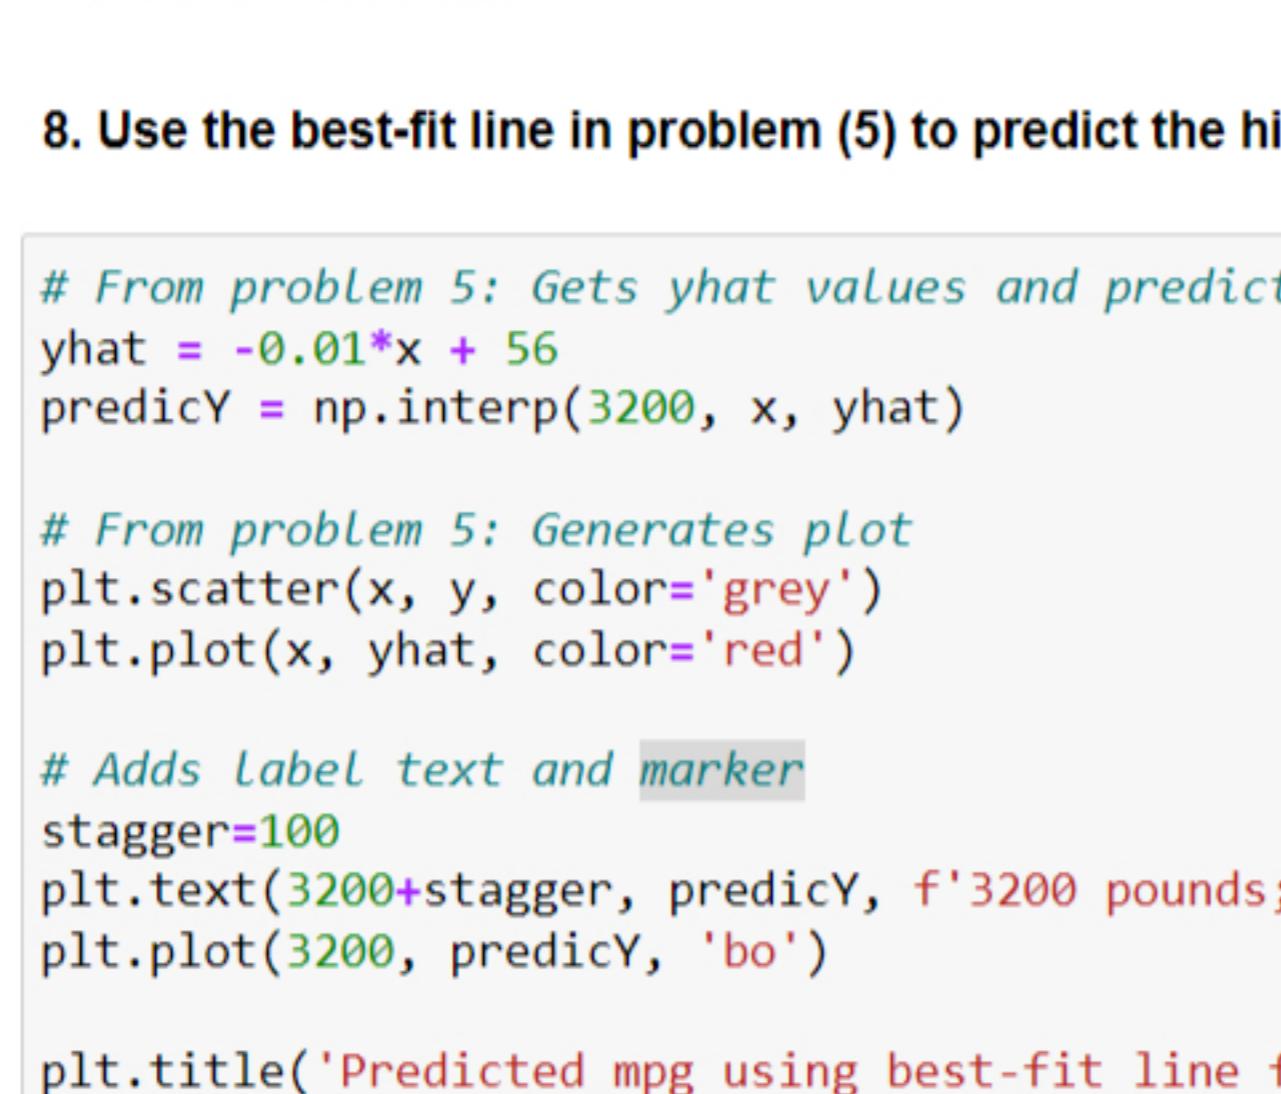
```
Out[31]: <matplotlib.lines.Line2D at 0x1a4403a0df0>
```



6. Use Python to find the best-fit line. The Scikit-learn package is a good choice to use for this.

```
In [33]: import numpy as np  
  
x, y = cars.weight, cars.hwy_mpg  
# Determines least squares polynomial fit for slope and y-intercept  
m, b = np.polyfit(x, y, 1)  
  
plt.scatter(x,y)  
  
# Adds regression line, m: -0.010547157168719583; b: 57.705174225744706  
plt.plot(x, m*x + b, color='red')
```

```
Out[33]: <matplotlib.lines.Line2D at 0x1a440429af0>
```



7. Find the root mean squared error (RMSE) of the prediction line you found in problem (4) and the actual best-fit line found in problem (5). How do these compare?

```
In [37]: def rmse(actual, predicted):  
    # Squares each element in List  
    def square(list):  
        return [i ** 2 for i in list]  
  
    # Equivalent to RMSE formula  
    sqrdiffs = square(predicted-actual)  
    return (1/len(actual)) * sum(sqrddiffs) ** 0.5
```

```
In [38]: # From problem 5:  
yhat = -0.01*x + 56
```

```
rmse(y, yhat)
```

```
Out[38]: 4.165945387977351
```

8. Use the best-fit line in problem (5) to predict the highway mpg of a car that weighs 3200 pounds.

```
In [63]: # From problem 5: Gets yhat values and predicts y  
yhat = -0.01*x + 56  
predicY = np.interp(3200, x, yhat)
```

```
# From problem 5: Generates plot  
plt.scatter(x, y, color='grey')
```

```
plt.plot(x, yhat, color='red')
```

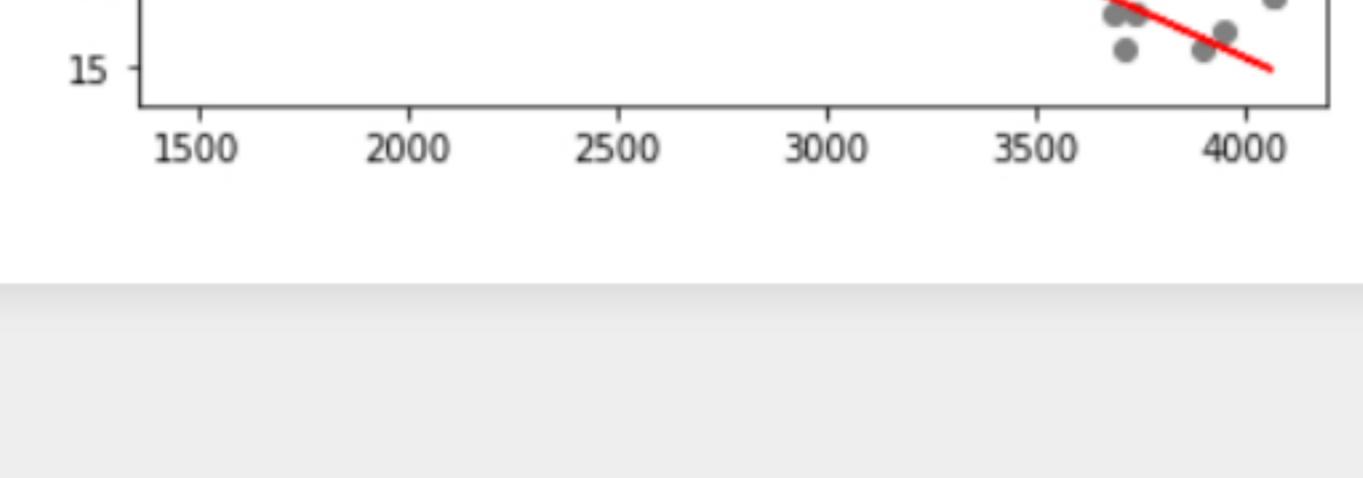
```
# Adds label text and marker  
stagger=100
```

```
plt.text(3200+stagger, predicY, f'3200 pounds; {predicY} mpg', color='blue')
```

```
plt.plot(3200, predicY, 'bo')
```

```
plt.title('Predicted mpg using best-fit line from #5')
```

```
Out[63]: Text(0.5, 1.0, 'Predicted mpg using best-fit line from #5')
```



```
In [65]: # From problem 6: Gets yhat values and predicts y  
m, b = np.polyfit(x, y, 1)
```

```
yhat = m*x + b
```

```
predicY = np.interp(3200, x, yhat)
```

```
# From problem 6: Generates plot  
plt.scatter(x,y, color='grey')
```

```
plt.plot(x, yhat, color='red')
```

```
# Adds label text and marker  
stagger=100
```

```
plt.text(3200+stagger, predicY, f'3200 pounds; {predicY} mpg', color='blue')
```

```
plt.plot(3200, predicY, 'bo')
```

```
plt.title('Predicted mpg using best-fit line from #6')
```

```
Out[65]: Text(0.5, 1.0, 'Predicted mpg using best-fit line from #6')
```

