

```
=====
Title: 3.2 Exercises
Author: Chad Wood
Date: 16 Dec 2021
Modified By: Chad Wood
Description: This program demonstrates the use of data wrangling techniques to better visualize data.
=====
```

US Population Growth

For this problem, you will be using the data set us pop data.csv. This data is from taken from the US Census has the US population every ten years from 1790 to 2010.

1. Import the data and create two new columns. Create one column that is the number of years since 1790. Create another column that is the population in millions.

```
In [1]: import pandas as pd
popUS = pd.read_csv('week3data/us_pop_data.csv')
```

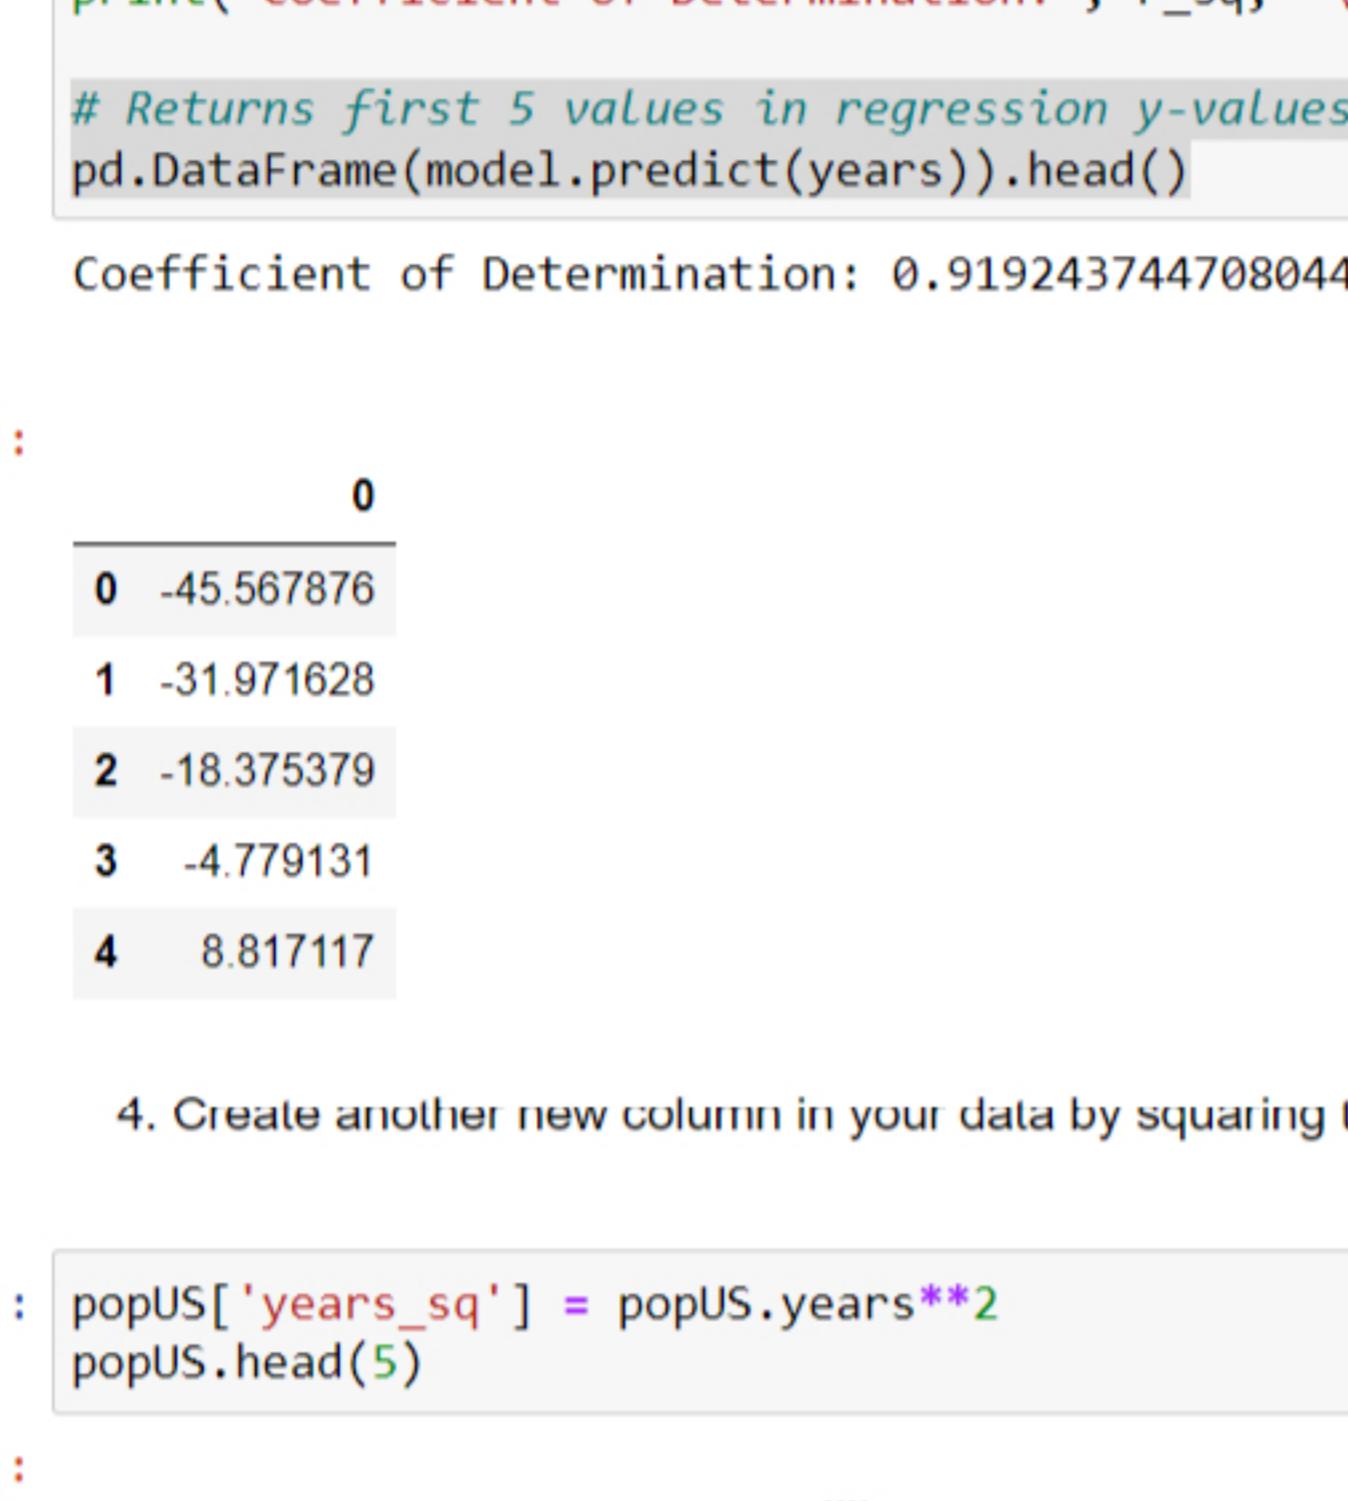
```
In [2]: popUS['years'] = popUS.year - 1790
popUS['pop_millions'] = popUS.us_pop / 1000000
popUS.head()
```

```
Out[2]:
   year  us_pop  years  pop_millions
0  1790  3929326      0     3.929326
1  1800  5308483     10     5.308483
2  1810  7239881     20     7.239881
3  1820  9638453     30     9.638453
4  1830 12866020     40    12.866020
```

2. Plot the US population (in millions) versus the years since 1790.

```
In [3]: import matplotlib.pyplot as plt
plt.plot(popUS.years, popUS.pop_millions)

plt.xlabel('Years Since 1790')
plt.ylabel('Pop in Millions')
plt.title('US Population Growth')
```



3. Create a linear regression model to predict the US population (in millions) to years from 1790. Find and report the R2-value of this model.

```
In [5]: import numpy as np
from sklearn.linear_model import LinearRegression
```

```
In [9]: pop = np.array(popUS.pop_millions)
```

```
In [76]: years = np.array(popUS.years).reshape((-1,1))

# Builds the Linear regression model
model = LinearRegression().fit(years, pop)

# Returns R2
r_sq = model.score(years, pop)
print('Coefficient of Determination:', r_sq, '\n')

# Returns first 5 values in regression y-values
pd.DataFrame(model.predict(years)).head()
```

Coefficient of Determination: 0.9192437447080442

```
Out[76]:
0
0  -45.567876
1  -31.971628
2  -18.375379
3  -4.779131
4  8.817117
```

4. Create another new column in your data by squaring the number of years since 1790.

```
In [14]: popUS['years_sq'] = popUS.years**2
popUS.head(5)
```

```
Out[14]:
   year  us_pop  years  pop_millions  years_sq
0  1790  3929326      0     3.929326      0
1  1800  5308483     10     5.308483    100
2  1810  7239881     20     7.239881    400
3  1820  9638453     30     9.638453    900
4  1830 12866020     40    12.866020   1600
```

5. Run another linear regression, where your input feature is the square of the number of years since 1790. Find and report the R2-value of this model.

```
In [77]: years_sq = np.array(popUS.years_sq).reshape((-1,1))
```

```
# Builds the Linear regression model
model = LinearRegression().fit(years_sq, pop)

# Returns R2
r_sq = model.score(years_sq, pop)
print('Coefficient of Determination:', r_sq, '\n')

# Returns first 5 values in regression y-values
pd.DataFrame(model.predict(years_sq)).head()
```

Coefficient of Determination: 0.9984915694986646

```
Out[77]:
0
0  1.360410
1  1.982413
2  3.648421
3  6.958435
4  11.312454
```

6. Plot the models you built on top of the data. Which one fits the data better? Is this apparent in your R2-values. Explain.

```
In [33]: # Create two subplots
f, (ax1, ax2) = plt.subplots(1, 2)

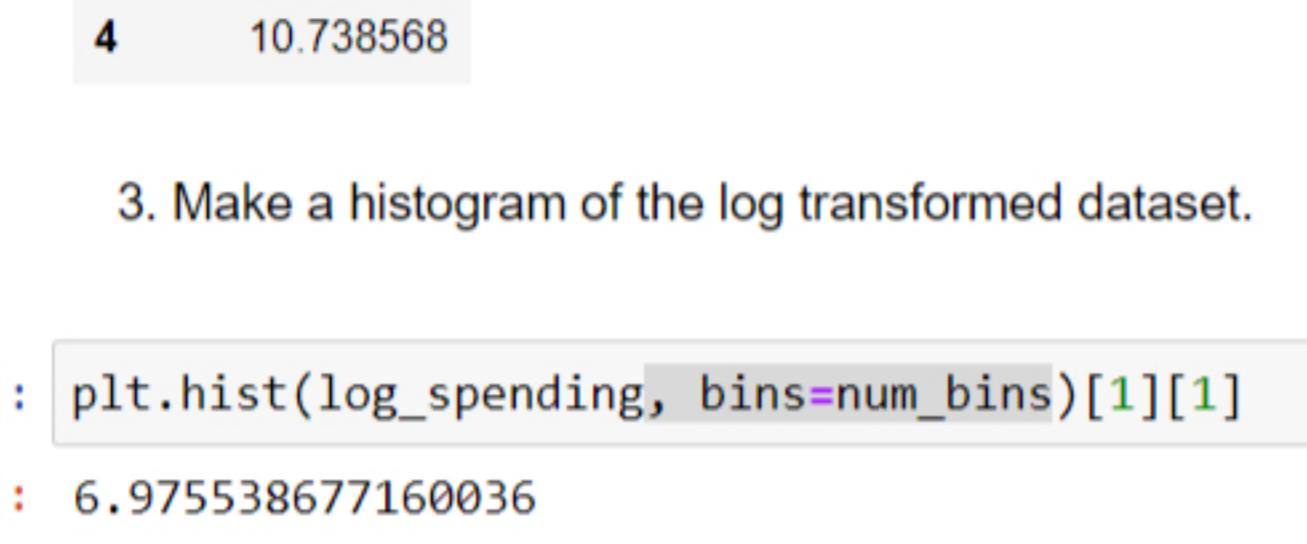
# Builds the Linear regression model; plots subplot
model = LinearRegression().fit(years, pop)

ax1.plot(years, model.predict(years), color='blue', label='Predicted')
ax1.plot(years, pop, color='grey', label='Actual')
ax1.legend()

# Builds the Linear regression model; plots subplot
model = LinearRegression().fit(years_sq, pop)

ax2.plot(years_sq, model.predict(years_sq), color='blue', label='Predicted; Years^2')
ax2.plot(years_sq, pop, color='grey', label='Actual')
ax2.legend()
```

```
Out[33]: <matplotlib.legend.Legend at 0x1f227dd0d0>
```



Linear regression lines are not good fits for exponential data, such as what's typically seen with population data. By squaring the years, we converted the years data from linear distribution to exponential distribution, and then built our regression model from that. This created a much better fit to our population data. This is also present in r^2 which represents a more perfect match the closer to 1 it is.

Customer Spending Data

For this problem, you will be using the data set customer spending.csv. This data set is modified version of the data from <https://archive.ics.uci.edu/ml/datasets/Wholesale+customers>.

1. Make a histogram of the customer spending amounts.

```
In [34]: spending = pd.read_csv('week3data/customer_spending.csv')
```

```
In [58]: import math
```

```
# Uses The Freedman-Diaconis rule to determine bin count
iqr = np.subtract(*np.percentile(spending, [75, 25]))
bin_width = 2 * iqr / len(spending)**(1/3)
num_bins = (spending.max() - spending.min()) / bin_width
```

```
# Rounds up
num_bins = math.ceil(num_bins)
```

```
In [68]: plt.hist(spending, bins=num_bins)[1]
```

```
Out[68]: 7122.34375
```


2. Make a new data set that is a log transformation of the spending amounts.

```
In [60]: log_spending = np.log(spending)
```

```
In [60]: log_spending.head()
```

```
Out[60]:
0  10.437405
1  10.112241
2  10.508077
3  10.217605
4  10.738568
```

3. Make a histogram of the log transformed dataset.

```
In [66]: plt.hist(log_spending, bins=num_bins)[1]
```

```
Out[66]: 6.975538677160036
```


4. Compare the two histograms. Discuss why it might be useful to apply a log transformation to this data for modeling purposes.

The original data contained a heavy positive skew in some cases, it's useful for better modelling. Since logarithmic calculations achieve a result closer to the standard curve without biasing the data.